




University of California  
San Francisco

# Introduction to Convolutional Neural Network


**Valentina Padoia, PhD**

**Musculoskeletal Quantitative Imaging Research Group  
Dept. of Radiology and Biomedical Imaging UCSF**



# Agenda

- **General concepts on Neural Networks**
  - Definitions, Architecture types, Backpropagation, Activation Functions
- **Multilayer ... Deep is better**
- **Handcraft Feature vs Feature Learning**
- **Convolutional Neural Networks**
  - Main Architectural Elements :
    - Convolutional
    - ReLu
    - Pooling
    - Fully Connected Layer



# Neural Network

- A Neural Network is an information processing paradigm that **is inspired by the biological nervous systems**, such as the human brain's information processing mechanism.
- **The key element of this paradigm is the structure of the information processing system.** It is composed of a large number of **highly interconnected processing elements** (neurons) working in unison to solve specific problems.
- NNs, “like people”, **learn by example**.
- A NN is configured for a specific application, such as pattern recognition or data classification, **through a learning process**.
- Learning in biological systems involves **adjustments to the synaptic connections** that exist between the neurons. **This is true of NNs as well.**





Figure 2 The anatomy of a neural network.



# Types of Neural Network Architecture

## Feed-forward networks

Feed-forward NNs allow signals to travel one way only; from input to output.

There is no feedback (loops) i.e. the output of any layer does not affect that same layer.

## Feedback networks

Feedback networks can have signals traveling in both directions by introducing loops in the network.


Feedback networks are dynamic; their 'state' is changing continuously until they reach an equilibrium point.

They remain at the equilibrium point until the input changes and a new equilibrium needs to be found.

Feedback architectures are also referred to as interactive or recurrent



# Learning Algorithm

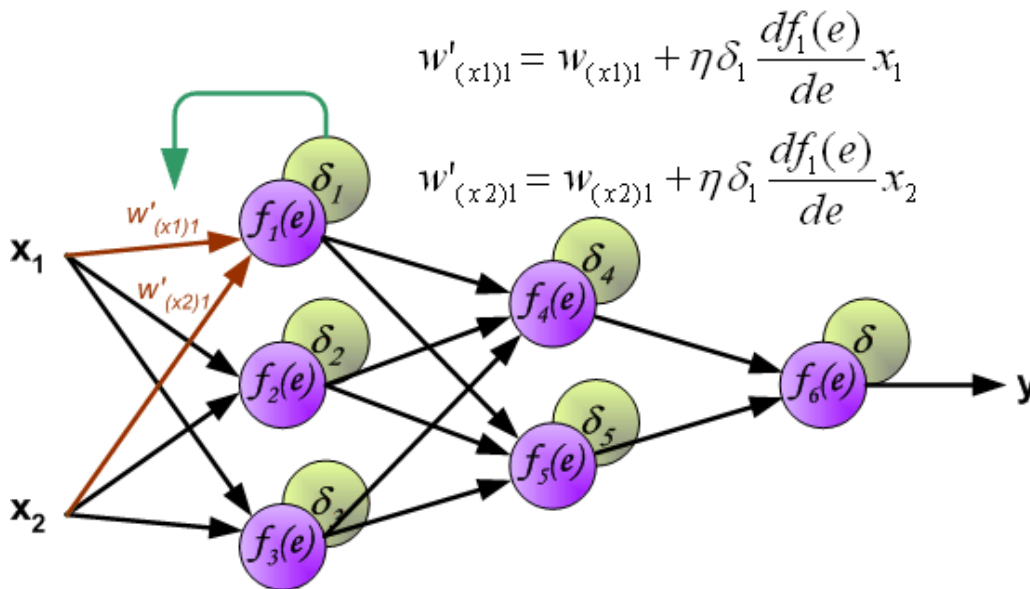



## Backpropagation:

Randomly initialize the parameters


Calculate total error at the top

Then calculate contributions to error, at each step going backwards






# The most simple neural network: Threshold Logic Unit




$$z = \begin{cases} 1 & \text{if } \sum_{i=1}^n x_i w_i \geq \theta \\ 0 & \text{if } \sum_{i=1}^n x_i w_i < \theta \end{cases}$$



# Training


$$\Delta w_i = c(t - z) x_i$$

- Where  $w_i$  is the weight from input  $i$  to perceptron node,
  - $c$  is the learning rate (the user decide this parameter),
  - $t_j$  is the target for the current instance,
  - $z$  is the current output, and
  - $x_i$  is  $i^{\text{th}}$  input
- 
- Least perturbation principle
    - Only change weights if there is an error
    - small  $c$  rather than changing weights sufficient to make current pattern correct
  - Iteratively apply a pattern from the training set and apply the perceptron rule
  - Each iteration through the training set is an *epoch*
  - Continue training until total training set error ceases to improve
  - Perceptron Convergence Theorem: Guaranteed to find a solution in finite time if a solution exists



# Let's try a simple example ....

1. Initialize the weights and the threshold




$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

$$1 \quad \text{if} \quad \sum_{i=1}^n x_i w_i \geq \theta$$

$z =$


$$0 \quad \text{if} \quad \sum_{i=1}^n x_i w_i < \theta$$





# Let's try a simple example ....

1. Initialize the weights and the threshold




	$x_1$	$x_2$	$t$
→	.8	.2	1
	.4	.3	0

$$1 \quad \text{if} \quad \sum_{i=1}^n x_i w_i \geq \theta$$


$z =$

$$0 \quad \text{if} \quad \sum_{i=1}^n x_i w_i < \theta$$



# Let's try a simple example ....

1. Initialize the weights and the threshold




$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

$$1 \quad \text{if} \quad \sum_{i=1}^n x_i w_i \geq \theta$$

$z =$

$$0 \quad \text{if} \quad \sum_{i=1}^n x_i w_i < \theta$$





# Let's try a simple example ....

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$






# Let's try a simple example ....

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$



	$x_1$	$x_2$	$t$
→	.8	.2	1
	.4	.3	0




# Let's try a simple example ....

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$



$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0






# Let's try a simple example ....

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$



$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

Training Set



$x_1$	$x_2$	$t$
.9	.2	1
.3	.3	0

Test Set




# Let's try a simple example ....

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$




$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

Training Set

$x_1$	$x_2$	$t$
.9	.2	1
.3	.3	0

Test Set






# Which pattern did we learn ?

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$




$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

Training Set

$x_1$	$x_2$	$t$
.9	.2	1
.3	.3	0

Test Set








# Which pattern did we learn ?

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$





$x_1$	$x_2$	$t$
.8	.2	1
.4	.3	0

Training Set

$x_1$	$x_2$	$t$
.9	.2	1
.3	.3	0

Test Set







## Add complexity: *input channels*

$$\Delta w_i = c(t - z) x_i$$

$$\Delta w_1 = 0.3 * (0 - 1) * 0.4 = -0.12$$

$$\Delta w_2 = 0.3 * (0 - 1) * 0.3 = -0.09$$




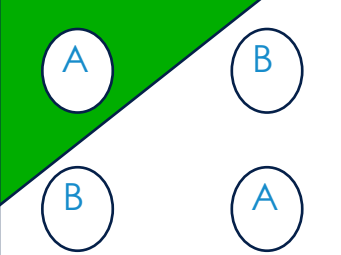
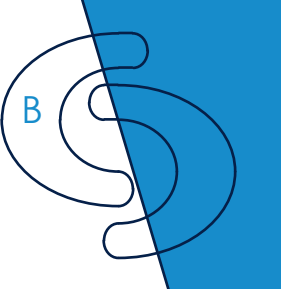

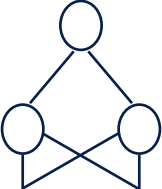
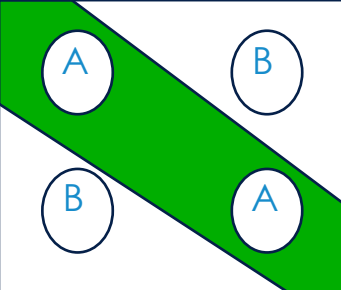
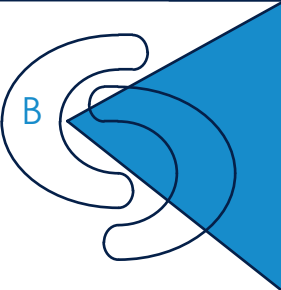
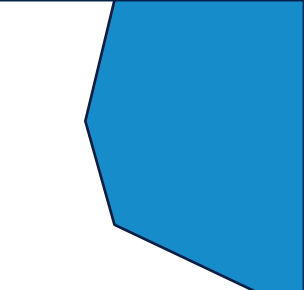
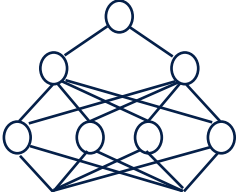
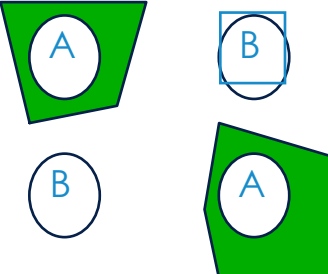
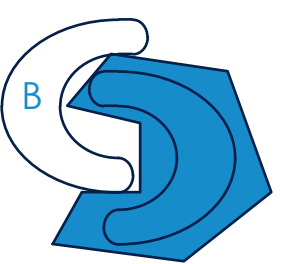
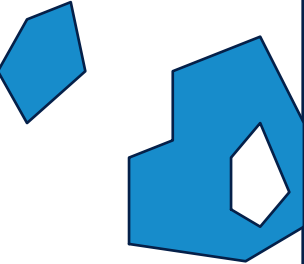


# Add complexity: *activation functions*

Activation function	Equation	Example	1D Graph
Unit step (Heaviside)	$\phi(z) = \begin{cases} 0, & z < 0, \\ 0.5, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Sign (Signum)	$\phi(z) = \begin{cases} -1, & z < 0, \\ 0, & z = 0, \\ 1, & z > 0, \end{cases}$	Perceptron variant	
Linear	$\phi(z) = z$	Adaline, linear regression	
Piece-wise linear	$\phi(z) = \begin{cases} 1, & z \geq \frac{1}{2}, \\ z + \frac{1}{2}, & -\frac{1}{2} < z < \frac{1}{2}, \\ 0, & z \leq -\frac{1}{2}, \end{cases}$	Support vector machine	
Logistic (sigmoid)	$\phi(z) = \frac{1}{1 + e^{-z}}$	Logistic regression, Multi-layer NN	
Hyperbolic tangent	$\phi(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$	Multi-layer NN	




# Multilayer Neural Network ...Why deep is better ...


Structure	Types of Decision Regions	Exclusive-OR Problem	Classes with Meshed regions	Most General Region Shapes
Single-Layer 	Half Plane Bounded By Hyper plane			
Two-Layer 	Convex Open Or Closed Regions			
Three-Layer 	Arbitrary (Complexity Limited by No. of Nodes)			



# Data coming from pixel in a image are not IID




...in computer vision/medical imaging



Classification


↑  
???




Maybe is not even 1 image but a stack of images or volumes ...



# Handcraft Feature vs Feature Learning



...in computer  
vision/medical  
imaging



Classification


## Feature Engineering

**Moments:** Average, std, entropy, kurtosis


**Domain Transformation:** FT, Wavelet, Spherical Harmonics

**Filtering:** Gabor Filters Fun Filters


Scale invariant feature transform **SIFT**




# Handcraft Feature vs Feature Learning




...in computer  
vision/medical  
imaging




Classification






# Handcraft Feature vs Feature Learning




...in computer  
vision/medical  
imaging

**Feature  
Learning**




Classification






# Handcraft Feature vs Feature Learning





...in computer  
vision/medical  
imaging

**Feature  
Learning**



Classification





# Convolutional Layers


- The core layer of CNNs
- The convolutional layer consists **of a set of filters**.
  - Each filter covers a spatially small portion of the input data.
- Each filter **is convolved across the dimensions of the input data**, producing a **multidimensional** feature map.
  - As we convolve the filter, we are computing the dot product between the parameters of the filter and the input.
- the network will learn filters that activate when they see some specific type of feature at some spatial position in the input.



# Convolutional Layer



Input



# Conv Hyperparameters

**Kernel Size:** controls how much **contextual information** is used to build feature maps

**Depth:** Depth corresponds to the **number of filters** we use for the convolution operation.


**Stride:** Stride is the number of pixels by which we slide our filter matrix over the input matrix. When the stride is 1 then we move the filters one pixel at a time. When the stride is 2, then the filters jump 2 pixels at a time as we slide them around. **Having a larger stride will produce smaller feature maps.**

**Zero-padding:** Sometimes, it is convenient to pad the input matrix with zeros around the border, so that we can apply the filter to bordering elements of our input image matrix. A nice feature of zero padding is that it allows us to control the size of the feature maps.

Adding zero-padding is also called ***wide convolution***, and not using zero-padding would be a ***narrow convolution***.




# CNN Architecture



...in computer  
vision/medical  
imaging




5x5x20  
Stride 2  
Pad 0




# Rectified Linear Unit

An additional operation called ReLU has been used after every Convolution operation ReLU stands for **R**ectified **L**inear **U**nit and is a non-linear operation

$$\text{Output} = \text{Max}(\text{zero}, \text{Input})$$

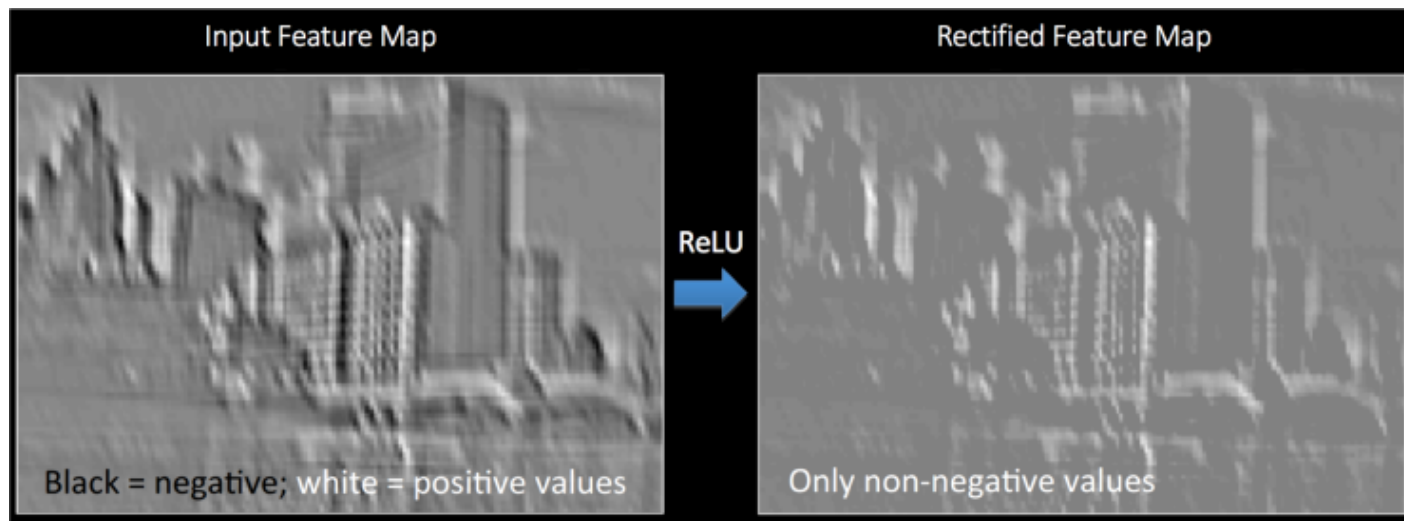





# Rectified Linear Unit


ReLU is an element wise operation (**applied per pixel**)

The purpose of ReLU is to **introduce non-linearity** in our ConvNet, since most of the real-world data we would want our ConvNet to learn would be non-linear (Convolution is a linear operation – element wise matrix multiplication and addition, so we account for non-linearity by introducing a non-linear function like ReLU).





# CNN Architecture




...in computer  
vision/medical  
imaging

ReLU




5x5x20  
Stride 2  
Pad 0






# Pooling/Downsampling


- Spatial Pooling (also called subsampling or downsampling) **reduces the dimensionality of each feature map** but retains the most important information.
- Spatial Pooling can be of different types: **Max, Average, Sum etc.**
- Can be learned
- It is not strictly necessary
- **Hyperparameters:**
  - Type, Filter Size, Stride






# Pooling/Downsampling

- Pooling makes the input representations (feature dimension) smaller and more manageable.
- Reduces the number of parameters and computations in the network, therefore, controlling overfitting
- Makes the network invariant to small transformations, distortions and translations in the input image (a small distortion in input will not change the output of Pooling – since we take the maximum / average value in a local neighborhood).



# CNN Architecture



...in computer  
vision/medical  
imaging

ReLU




Basic building blocks of any CNN


5x5x20  
Stride 2  
Pad 0

Max  
Pooling  
2X2  
Stride 2


Convolutional Layer




# CNN Architecture




...in computer  
vision/medical  
imaging




Several convolutional layers are used to  
build a **Deep** CNN with different hyper  
parameters



# CNN Architecture



...in computer  
vision/medical  
imaging




Several convolutional layers are used to  
build a **Deep** CNN with different hyper  
parameters


**Feature Learning**




# CNN Architecture



...in computer vision/medical imaging




Several convolutional layers are used to build a **Deep** CNN with different hyper parameters



Fully Connected Layer

**Feature Learning**


**Classifier**




The Fully Connected layer is a traditional **Multi Layer Perceptron** that uses a **Softmax** activation function in the output layer




# LeNet





# AlexNet

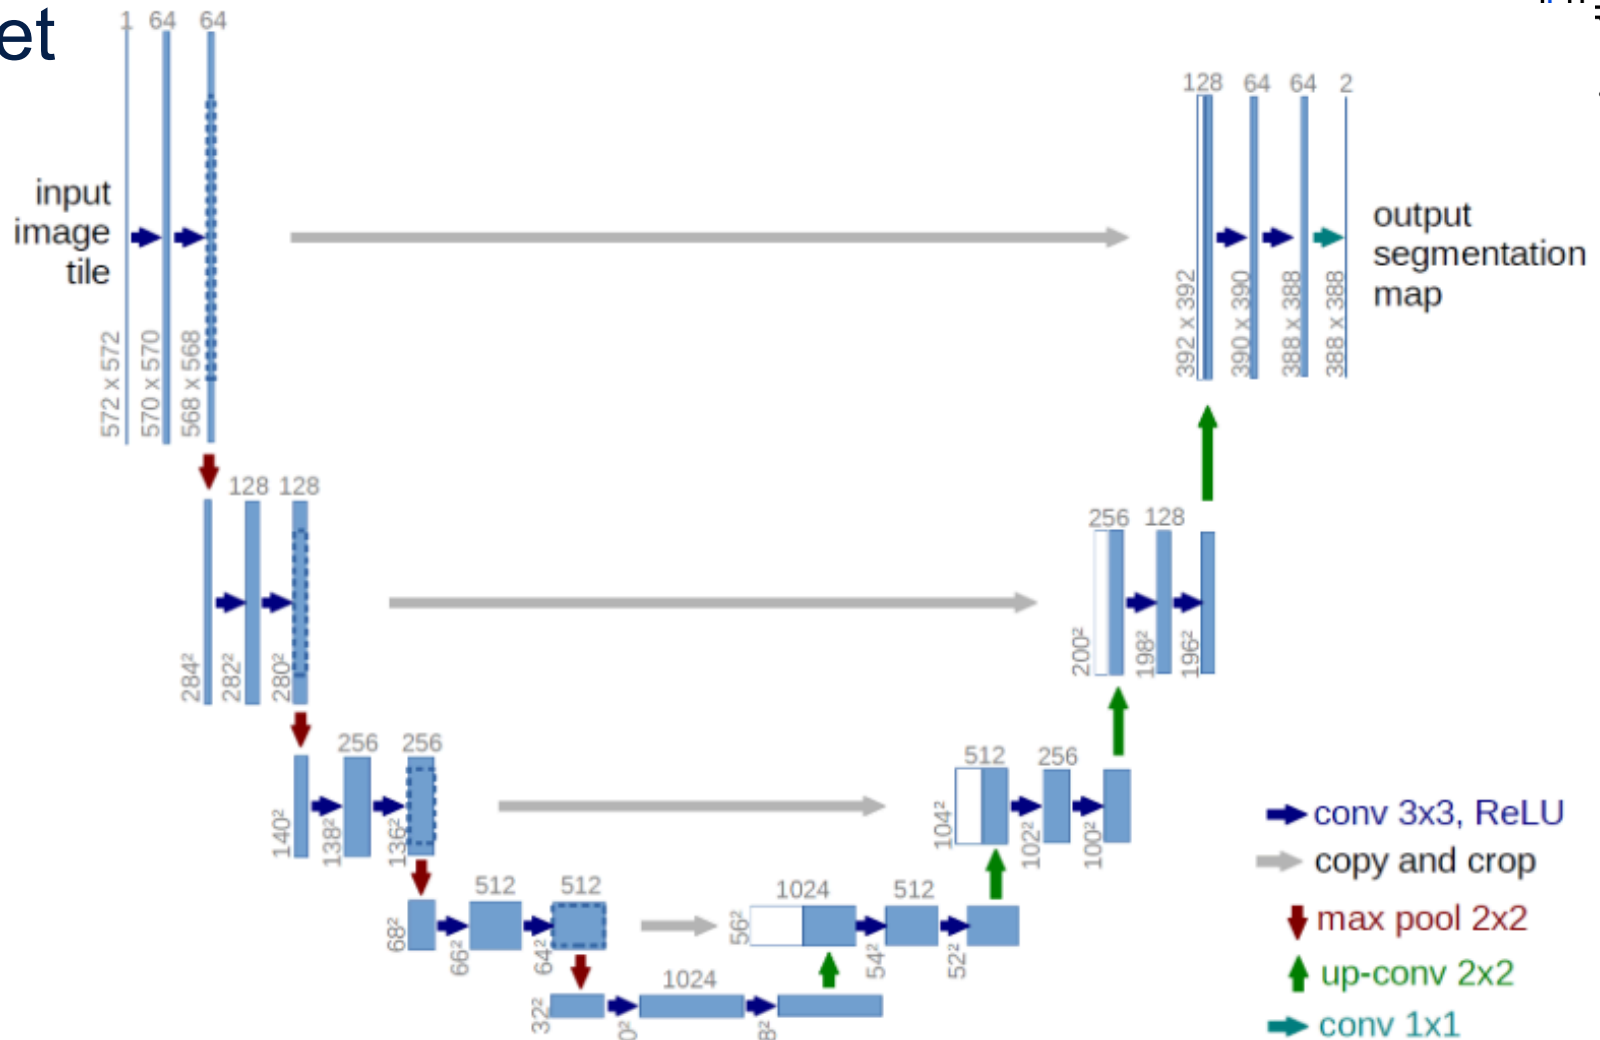




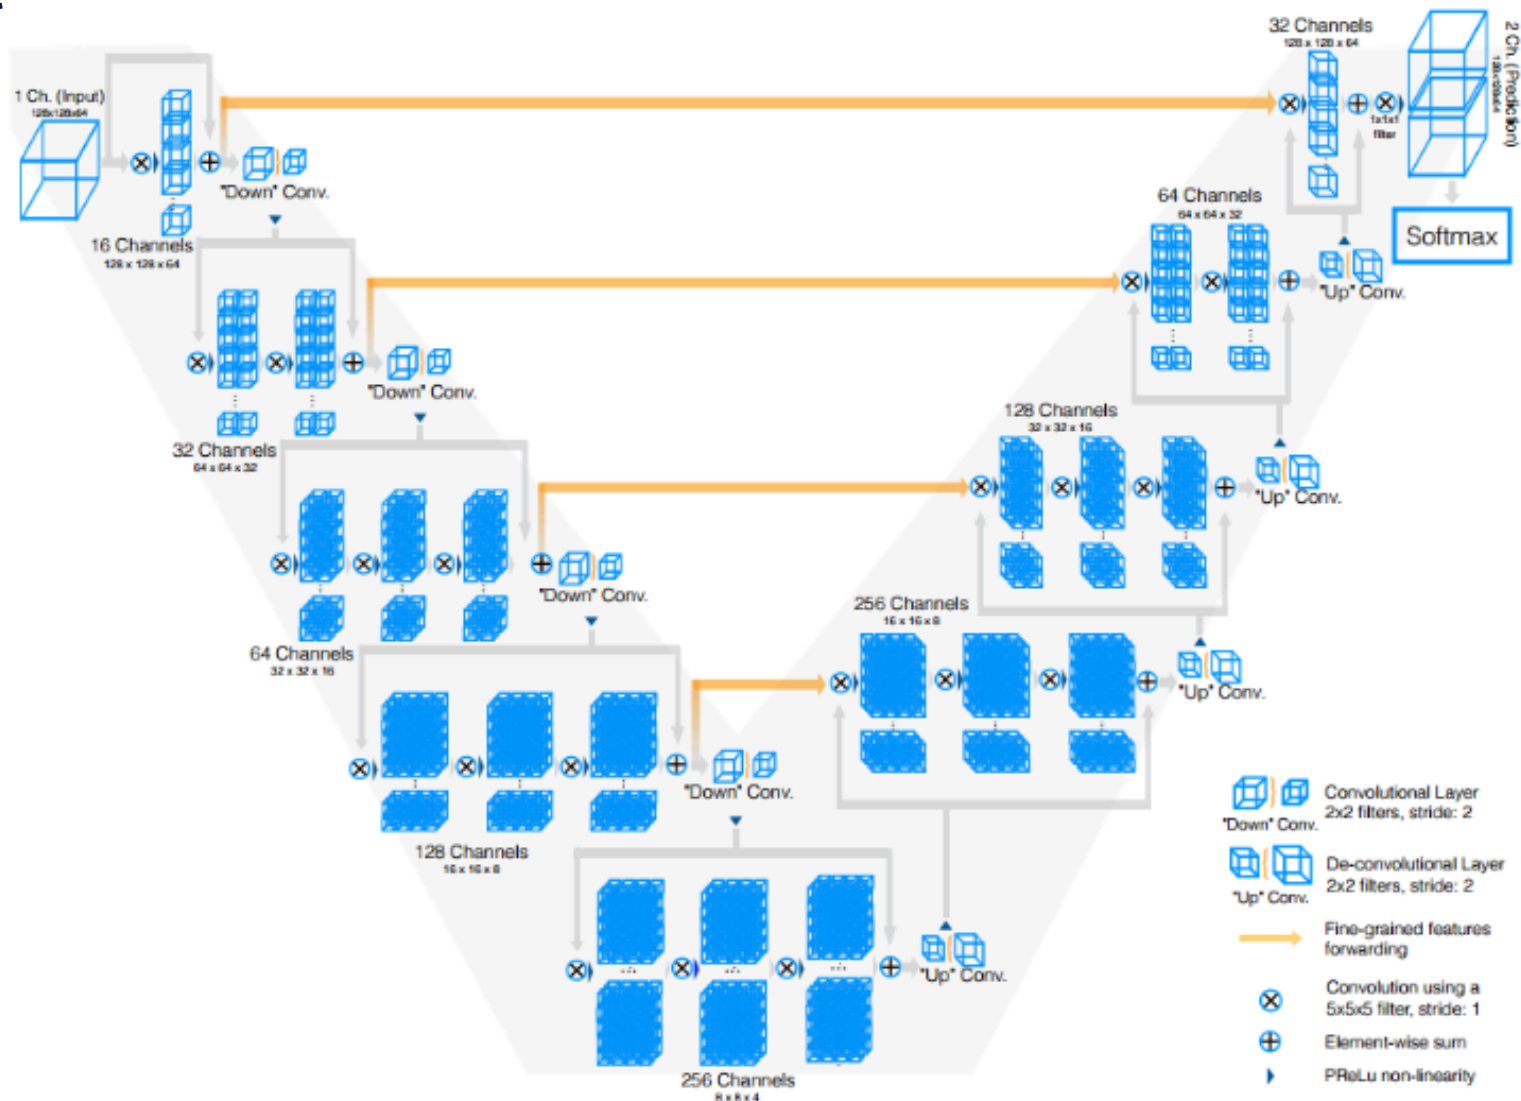





# U-net



# V-net





# Generative adversarial network (**GAN**)

